

Introduction to Machine Learning

Maximum Margin Methods

Varun Chandola

Computer Science & Engineering
State University of New York at Buffalo
Buffalo, NY, USA
chandola@buffalo.edu



Maximum Margin Classifiers

- Linear Classification via Hyperplanes

- Concept of Margin

Support Vector Machines

- SVM Learning

- Solving SVM Optimization Problem

Constrained Optimization and Lagrange Multipliers

- Karun-Kuhn-Tucker Conditions


- Support Vectors

- Optimization Constraints

Maximum Margin Classifiers

$$y = \mathbf{w}^T \mathbf{x} + b$$


- ▶ Remember the Perceptron!
- ▶ If data is linearly separable
 - ▶ Perceptron training guarantees learning the decision boundary
- ▶ There can be other boundaries
 - ▶ Depends on initial value for \mathbf{w}



Maximum Margin Classifiers


$$y = \mathbf{w}^T \mathbf{x} + b$$

- ▶ Remember the Perceptron!
- ▶ If data is linearly separable
 - ▶ Perceptron training guarantees learning the decision boundary
- ▶ There can be other boundaries
 - ▶ Depends on initial value for \mathbf{w}
- ▶ **But what is the best boundary?**



Linear Hyperplane

- ▶ Separates a D -dimensional space into two half-spaces
- ▶ Defined by $\mathbf{w} \in \mathbb{R}^D$
 - ▶ *Orthogonal* to the hyperplane
 - ▶ This \mathbf{w} goes through the origin
 - ▶ How do you check if a point lies “above” or “below” \mathbf{w} ?
 - ▶ What happens for points **on** \mathbf{w} ?



Make hyperplane not go through origin

- ▶ Add a bias b
 - ▶ $b > 0$ - move along \mathbf{w}
 - ▶ $b < 0$ - move opposite to \mathbf{w}
- ▶ How to check if point lies above or below \mathbf{w} ?
 - ▶ If $\mathbf{w}^\top \mathbf{x} + b > 0$ then \mathbf{x} is *above*
 - ▶ Else, *below*


Line as a Decision Surface

- ▶ Decision boundary represented by the hyperplane \mathbf{w}
- ▶ For binary classification, \mathbf{w} points **towards** the positive class

Decision Rule


$$y = \text{sign}(\mathbf{w}^\top \mathbf{x} + b)$$

- ▶ $\mathbf{w}^\top \mathbf{x} + b > 0 \Rightarrow y = +1$
- ▶ $\mathbf{w}^\top \mathbf{x} + b < 0 \Rightarrow y = -1$



What is Best Hyperplane Separator

- ▶ **Perceptron** can find a hyperplane that separates the data
 - ▶ ... if the data is linearly separable
- ▶ But there can be many choices!
- ▶ Find the one with best separability (largest margin)
- ▶ Gives better generalization performance
 1. Intuitive reason
 2. Theoretical foundations



What is a Margin?

- ▶ **Margin** is the distance between an example and the decision line
- ▶ Denoted by γ
- ▶ For a positive point:

$$\gamma = \frac{\mathbf{w}^\top \mathbf{x} + b}{\|\mathbf{w}\|}$$


- ▶ For a negative point:

$$\gamma = -\frac{\mathbf{w}^\top \mathbf{x} + b}{\|\mathbf{w}\|}$$

Functional Interpretation

- ▶ Margin **positive** if prediction is **correct**; **negative** if prediction is **incorrect**

Maximum Margin Principle



Support Vector Machines

- ▶ A hyperplane based classifier defined by \mathbf{w} and b
- ▶ Like perceptron
- ▶ Find hyperplane with *maximum separation margin* on the training data
- ▶ Assume that data is linearly separable (will relax this later)
 - ▶ Zero training error (loss)

SVM Prediction Rule

$$y = \text{sign}(\mathbf{w}^\top \mathbf{x} + b)$$

SVM Learning

- ▶ **Input:** Training data $\{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)\}$
- ▶ **Objective:** Learn \mathbf{w} and b that maximizes the margin

- ▶ SVM learning task as an optimization problem
- ▶ Find \mathbf{w} and b that gives zero training error
- ▶ Maximizes the margin ($= \frac{2}{\|\mathbf{w}\|}$)
- ▶ Same as minimizing $\|\mathbf{w}\|$

Optimization Formulation

$$\begin{aligned} & \underset{\mathbf{w}, b}{\text{minimize}} && \frac{\|\mathbf{w}\|^2}{2} \\ & \text{subject to} && y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1, \quad i = 1, \dots, n. \end{aligned}$$

- ▶ **Optimization** with N linear inequality constraint

A Different Interpretation of Margin

- ▶ What impact does the margin have on \mathbf{w} ?

A Different Interpretation of Margin

- ▶ What impact does the margin have on \mathbf{w} ?
- ▶ Large margin \Rightarrow Small $\|\mathbf{w}\|$

A Different Interpretation of Margin

- ▶ What impact does the margin have on \mathbf{w} ?
- ▶ Large margin \Rightarrow Small $\|\mathbf{w}\|$
- ▶ Small $\|\mathbf{w}\| \Rightarrow$ regularized/simple solutions

A Different Interpretation of Margin

- ▶ What impact does the margin have on \mathbf{w} ?
- ▶ Large margin \Rightarrow Small $\|\mathbf{w}\|$
- ▶ Small $\|\mathbf{w}\| \Rightarrow$ regularized/simple solutions
- ▶ Simple solutions \Rightarrow Better generalizability (*Occam's Razor*)

A Different Interpretation of Margin

- ▶ What impact does the margin have on \mathbf{w} ?
- ▶ Large margin \Rightarrow Small $\|\mathbf{w}\|$
- ▶ Small $\|\mathbf{w}\| \Rightarrow$ regularized/simple solutions
- ▶ Simple solutions \Rightarrow Better generalizability (*Occam's Razor*)
- ▶ Computational Learning Theory provides a formal justification [1]

Solving the Optimization Problem

Optimization Formulation

$$\begin{aligned} & \underset{\mathbf{w}, b}{\text{minimize}} && \frac{\|\mathbf{w}\|^2}{2} \\ & \text{subject to} && y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1, \quad i = 1, \dots, n. \end{aligned}$$

- ▶ There is an quadratic objective function to minimize with N inequality constraints
- ▶ “Off-the-shelf” packages - quadprog (MATLAB), CVXOPT
- ▶ Is that the best way?

Basic Optimization

$$\underset{x,y}{\text{minimize}} \quad f(x,y) = x^2 + 2y^2 - 2$$

Basic Optimization

$$\underset{x,y}{\text{minimize}} \quad f(x,y) = x^2 + 2y^2 - 2$$

$$\begin{aligned} &\underset{x,y}{\text{minimize}} \quad f(x,y) = x^2 + 2y^2 - 2 \\ &\text{subject to} \quad h(x,y) = x + y - 1 = 0. \end{aligned}$$

Lagrange Multipliers - A Primer

- ▶ Tool for solving constrained optimization problems of differentiable functions

$$\begin{array}{ll}\text{minimize}_{x,y} & f(x,y) = x^2 + 2y^2 - 2 \\ \text{subject to} & h(x,y) : x + y - 1 = 0.\end{array}$$

- ▶ A Lagrangian multiplier (β) lets you combine the two equations into one

Lagrange Multipliers - A Primer

- ▶ Tool for solving constrained optimization problems of differentiable functions

$$\begin{array}{ll}\text{minimize}_{x,y} & f(x,y) = x^2 + 2y^2 - 2 \\ \text{subject to} & h(x,y) : x + y - 1 = 0.\end{array}$$

- ▶ A Lagrangian multiplier (β) lets you combine the two equations into one

$$\text{minimize}_{x,y,\beta} \quad L(x,y,\beta) = f(x,y) + \beta h(x,y)$$

Multiple Constraints

$$\begin{array}{ll} \underset{x,y,z}{\text{minimize}} & f(x,y,z) = x^2 + 4y^2 + 2z^2 + 6y + z \\ \text{subject to} & h_1(x,y,z) : \quad \quad \quad x + z^2 - 1 = 0 \\ & h_2(x,y,z) : \quad \quad \quad x^2 + y^2 - 1 = 0. \end{array}$$

Multiple Constraints

$$\begin{array}{ll}\text{minimize}_{x,y,z} & f(x,y,z) = x^2 + 4y^2 + 2z^2 + 6y + z \\ \text{subject to} & h_1(x,y,z) : \quad x + z^2 - 1 = 0 \\ & h_2(x,y,z) : \quad x^2 + y^2 - 1 = 0.\end{array}$$

$$L(x,y,z,\beta) = f(x,y,z) + \sum_i \beta_i h_i(x,y,z)$$

Handling Inequality Constraints

$$\begin{array}{ll} \underset{x,y}{\text{minimize}} & f(x,y) = x^3 + y^2 \\ \text{subject to} & g(x) : x^2 - 1 \leq 0. \end{array}$$

Handling Inequality Constraints

$$\begin{array}{ll} \underset{x,y}{\text{minimize}} & f(x,y) = x^3 + y^2 \\ \text{subject to} & g(x) : x^2 - 1 \leq 0. \end{array}$$

- ▶ Inequality constraints are **transferred** as constraints on the Lagrangian, α

Handling Both Types of Constraints

$$\begin{array}{ll} \underset{\mathbf{w}}{\text{minimize}} & f(\mathbf{w}) \\ \text{subject to} & g_i(\mathbf{w}) \leq 0 \quad i = 1, \dots, k \\ \text{and} & h_i(\mathbf{w}) = 0 \quad i = 1, \dots, l. \end{array}$$

Generalized Lagrangian

$$L(\mathbf{w}, \alpha, \beta) = f(\mathbf{w}) + \sum_{i=1}^k \alpha_i g_i(\mathbf{w}) + \sum_{i=1}^l \beta_i h_i(\mathbf{w})$$

subject to, $\alpha_i \geq 0, \forall i$

Primal Optimization

- ▶ Let θ_P be defined as:

$$\theta_P(\mathbf{w}) = \max_{\alpha, \beta: \alpha_i \geq 0} L(\mathbf{w}, \alpha, \beta)$$

- ▶ One can prove that the optimal value for the original constrained problem is same as:

$$p^* = \min_{\mathbf{w}} \theta_P(\mathbf{w}) = \min_{\mathbf{w}} \max_{\alpha, \beta: \alpha_i \geq 0} L(\mathbf{w}, \alpha, \beta)$$

Primal and Dual Formulations (II)

Dual Optimization

- ▶ Consider θ_D , defined as:

$$\theta_D(\alpha, \beta) = \min_{\mathbf{w}} L(\mathbf{w}, \alpha, \beta)$$

- ▶ The **dual** optimization problem can be posed as:

$$d^* = \max_{\alpha, \beta: \alpha_i \geq 0} \theta_D(\alpha, \beta) = \max_{\alpha, \beta: \alpha_i \geq 0} \min_{\mathbf{w}} L(\mathbf{w}, \alpha, \beta)$$

$$d^* == p^*?$$

- ▶ Note that $d^* \leq p^*$
- ▶ “Max min” of a function is always less than or equal to “Min max”
- ▶ When will they be equal?
 - ▶ $f(\mathbf{w})$ is convex
 - ▶ Constraints are affine

Relation between primal and dual

- ▶ In general $d^* \leq p^*$, for SVM optimization the equality holds
- ▶ Certain conditions should be true
- ▶ Known as the **Karun-Kuhn-Tucker** conditions
- ▶ For $d^* = p^* = L(\mathbf{w}^*, \alpha^*, \beta^*)$:

$$\begin{aligned}\frac{\partial}{\partial \mathbf{w}} L(\mathbf{w}^*, \alpha^*, \beta^*) &= 0 \\ \frac{\partial}{\partial \beta_i} L(\mathbf{w}^*, \alpha^*, \beta^*) &= 0, \quad i = 1, \dots, l \\ \alpha_i^* g_i(\mathbf{w}^*) &= 0, \quad i = 1, \dots, k \\ g_i(\mathbf{w}^*) &\leq 0, \quad i = 1, \dots, k \\ \alpha_i^* &\geq 0, \quad i = 1, \dots, k\end{aligned}$$

Lagrangian Multipliers for SVM

Optimization Formulation

$$\begin{aligned} & \underset{\mathbf{w}, b}{\text{minimize}} && \frac{\|\mathbf{w}\|^2}{2} \\ & \text{subject to} && y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1, \quad i = 1, \dots, n. \end{aligned}$$

Lagrangian Multipliers for SVM

Optimization Formulation

$$\begin{aligned} & \underset{\mathbf{w}, b}{\text{minimize}} && \frac{\|\mathbf{w}\|^2}{2} \\ & \text{subject to} && y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1, \quad i = 1, \dots, n. \end{aligned}$$

A Toy Example

- ▶ $\mathbf{x} \in \Re^2$
- ▶ Two training points:
$$\mathbf{x}_1, y_1 = (1, 1), -1$$
$$\mathbf{x}_2, y_2 = (2, 2), +1$$
- ▶ Find the best hyperplane $\mathbf{w} = (w_1, w_2)$

Optimization problem for the toy example

$$\begin{array}{ll}\underset{\mathbf{w}}{\text{minimize}} & f(\mathbf{w}) = \frac{1}{2} \|\mathbf{w}\|^2 \\ \text{subject to} & g_1(\mathbf{w}, b) = y_1(\mathbf{w}^\top \mathbf{x}_1 + b) - 1 \geq 0 \\ & g_2(\mathbf{w}, b) = y_2(\mathbf{w}^\top \mathbf{x}_2 + b) - 1 \geq 0.\end{array}$$

Optimization problem for the toy example

$$\begin{array}{ll}\underset{\mathbf{w}}{\text{minimize}} & f(\mathbf{w}) = \frac{1}{2} \|\mathbf{w}\|^2 \\ \text{subject to} & g_1(\mathbf{w}, b) = y_1(\mathbf{w}^\top \mathbf{x}_1 + b) - 1 \geq 0 \\ & g_2(\mathbf{w}, b) = y_2(\mathbf{w}^\top \mathbf{x}_2 + b) - 1 \geq 0.\end{array}$$

- Substituting actual values for \mathbf{x}_1, y_1 and \mathbf{x}_2, y_2 .

$$\begin{array}{ll}\underset{\mathbf{w}}{\text{minimize}} & f(\mathbf{w}) = \frac{1}{2} \|\mathbf{w}\|^2 \\ \text{subject to} & g_1(\mathbf{w}, b) = -(\mathbf{w}^\top \mathbf{x}_1 + b) - 1 \geq 0 \\ & g_2(\mathbf{w}, b) = (\mathbf{w}^\top \mathbf{x}_2 + b) - 1 \geq 0.\end{array}$$

Optimization Formulation

$$\begin{aligned} & \underset{\mathbf{w}, b}{\text{minimize}} && \frac{\|\mathbf{w}\|^2}{2} \\ & \text{subject to} && y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1, \quad i = 1, \dots, n. \end{aligned}$$

- Introducing [Lagrange Multipliers](#), $\alpha_i, i = 1, \dots, n$

Rewriting as a (primal) Lagrangian

$$\begin{aligned} & \underset{\mathbf{w}, b, \alpha}{\text{minimize}} && L_P(\mathbf{w}, b, \alpha) = \frac{\|\mathbf{w}\|^2}{2} + \sum_{i=1}^n \alpha_i \{1 - y_i(\mathbf{w}^\top \mathbf{x}_i + b)\} \\ & \text{subject to} && \alpha_i \geq 0 \quad i = 1, \dots, n. \end{aligned}$$

Solving the Lagrangian

- ▶ Set gradient of L_P to 0

$$\frac{\partial L_P}{\partial \mathbf{w}} = 0 \Rightarrow \mathbf{w} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i$$

$$\frac{\partial L_P}{\partial b} = 0 \Rightarrow \sum_{i=1}^n \alpha_i y_i = 0$$

- ▶ Substituting in L_P to get the dual L_D

Solving the Lagrangian

- Set gradient of L_P to 0

$$\frac{\partial L_P}{\partial \mathbf{w}} = 0 \Rightarrow \mathbf{w} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i$$

$$\frac{\partial L_P}{\partial b} = 0 \Rightarrow \sum_{i=1}^n \alpha_i y_i = 0$$

- Substituting in L_P to get the dual L_D

Dual Lagrangian Formulation

$$\begin{aligned} \underset{b, \alpha}{\text{maximize}} \quad & L_D(\mathbf{w}, b, \alpha) = \sum_{i=1}^n \alpha_i - \frac{1}{2} \sum_{m, n=1}^N \alpha_m \alpha_n y_m y_n (\mathbf{x}_m^\top \mathbf{x}_n) \\ \text{subject to} \quad & \sum_{i=1}^n \alpha_i y_i = 0, \alpha_i \geq 0 \quad i = 1, \dots, n. \end{aligned}$$

Solving the Dual

- ▶ Dual Lagrangian is a *quadratic programming problem* in α_i 's
 - ▶ Use “off-the-shelf” solvers
- ▶ Having found α_i 's

$$\mathbf{w} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i$$

- ▶ What will be the bias term b ?

Investigating Karush Kuhn Tucker Conditions

- ▶ For the primal and dual formulations
- ▶ We can optimize the dual formulation (as shown earlier)
- ▶ Solution should satisfy the **Karush-Kuhn-Tucker** (KKT) Conditions

The Kahrn-Kuhn-Tucker Conditions

$$\frac{\partial}{\partial \mathbf{w}} L_P(\mathbf{w}, b, \alpha) = \mathbf{w} - \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i = 0 \quad (1)$$

$$\frac{\partial}{\partial b} L_P(\mathbf{w}, b, \alpha) = - \sum_{i=1}^n \alpha_i y_i = 0 \quad (2)$$

$$1 - y_i \{\mathbf{w}^\top \mathbf{x}_i + b\} \leq 0 \quad (3)$$

$$\alpha_i \geq 0 \quad (4)$$

$$\alpha_i (1 - y_i \{\mathbf{w}^\top \mathbf{x}_i + b\}) = 0 \quad (5)$$

Estimating Bias b

- ▶ Use KKT condition #5
- ▶ For $\alpha_i > 0$

$$(y_i\{\mathbf{w}^\top \mathbf{x}_i + b\} - 1) = 0$$

- ▶ Which means that:

$$b = -\frac{\max_{n:y_i=-1} \mathbf{w}^\top \mathbf{x}_i + \min_{n:y_i=1} \mathbf{w}^\top \mathbf{x}_i}{2}$$

Key Observation from Dual Formulation

Most α_i 's are 0

- ▶ KKT condition #5:


$$\alpha_i(1 - y_i\{\mathbf{w}^\top \mathbf{x}_i + b\}) = 0$$

- ▶ If \mathbf{x}_i **not** on margin

$$y_i\{\mathbf{w}^\top \mathbf{x}_i + b\} > 1$$


$$\Rightarrow \alpha_i = 0$$

- ▶ $\alpha_i \neq 0$ only for \mathbf{x}_i on margin
- ▶ These are the **support vectors**
- ▶ Only need these for prediction



What have we seen so far?

- ▶ For linearly separable data, SVM learns a weight vector \mathbf{w}
- ▶ Maximizes the margin
- ▶ SVM training is a **constrained optimization problem**
 - ▶ Each training example should lie outside the margin
 - ▶ N constraints



What if data is not linearly separable?

- ▶ Cannot go for zero training error
- ▶ Still learn a maximum margin hyperplane


What if data is not linearly separable?

- ▶ Cannot go for zero training error
- ▶ Still learn a maximum margin hyperplane
 1. Allow some examples to be misclassified
 2. Allow some examples to fall **inside** the margin

What if data is not linearly separable?

- ▶ Cannot go for zero training error
- ▶ Still learn a maximum margin hyperplane
 1. Allow some examples to be misclassified
 2. Allow some examples to fall **inside** the margin
- ▶ How do you set up the optimization for SVM training

Cutting Some Slack



Introducing Slack Variables

- **Separable Case:** To ensure zero training loss, constraint was

$$y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 \quad \forall i = 1 \dots n$$

Introducing Slack Variables

- ▶ **Separable Case:** To ensure zero training loss, constraint was

$$y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 \quad \forall i = 1 \dots n$$

- ▶ **Non-separable Case:** Relax the constraint

$$y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 - \xi_i \quad \forall i = 1 \dots n$$

- ▶ ξ_i is called **slack variable** ($\xi_i \geq 0$)
- ▶ For misclassification, $\xi_i > 1$

Relaxing the Constraint

- ▶ It is OK to have some misclassified training examples
 - ▶ Some ξ_i 's will be non-zero

Relaxing the Constraint

- ▶ It is OK to have some misclassified training examples
 - ▶ Some ξ_i 's will be non-zero
- ▶ Minimize the number of such examples

- ▶ Minimize $\sum_{i=1}^n \xi_i$

- ▶ Optimization Problem for Non-Separable Case

$$\begin{aligned} \underset{\mathbf{w}, b}{\text{minimize}} \quad & f(\mathbf{w}, b) = \|\mathbf{w}\|^2 + C \sum_{i=1}^n \xi_i \\ \text{subject to} \quad & y_i(\mathbf{w}^\top \mathbf{x}_i + b) \geq 1 - \xi_i, \xi_i \geq 0 \quad i = 1, \dots, n. \end{aligned}$$

- ▶ C controls the impact of margin and the margin error.

Estimating Weights

- ▶ What is the role of C ?
- ▶ Similar optimization procedure as for the separable case (QP for the dual)
- ▶ Weights have the same expression

$$\mathbf{w} = \sum_{i=1}^n \alpha_i y_i \mathbf{x}_i$$

- ▶ Support vectors are slightly different
 1. Points on the margin ($\xi_i = 0$)
 2. Inside the margin but on the correct side ($0 < \xi_i < 1$)
 3. On the wrong side of the hyperplane ($\xi_i \geq 1$)

Concluding Remarks on SVM

- ▶ Training time for SVM training is $O(n^3)$
- ▶ Many *faster* but approximate approaches exist
 - ▶ Approximate QP solvers
 - ▶ Online training
- ▶ SVMs can be extended in different ways
 1. Non-linear boundaries (**kernel trick**)
 2. Multi-class classification
 3. Probabilistic output
 4. Regression (Support Vector Regression)

References



V. Vapnik.

Statistical learning theory.

Wiley, 1998.