Introduction to Machine Learning

Introduction

Varun Chandola

Computer Science & Engineering
State University of New York at Buffalo
Buffalo, NY, USA
chandola@buffalo.edu

University at Buffalo
Department of Computer Science and Engineering
School of Engeering and Applied Sciences

Outline

Making Machines Intelligent

The world is woven from billions of lives, every strand crossing every other. What we call premonition is just movement of the web. If you could attenuate to every strand of quivering data, the future would be entirely calculable, as inevitable as mathematics.

Sherlock Holmes, 2017

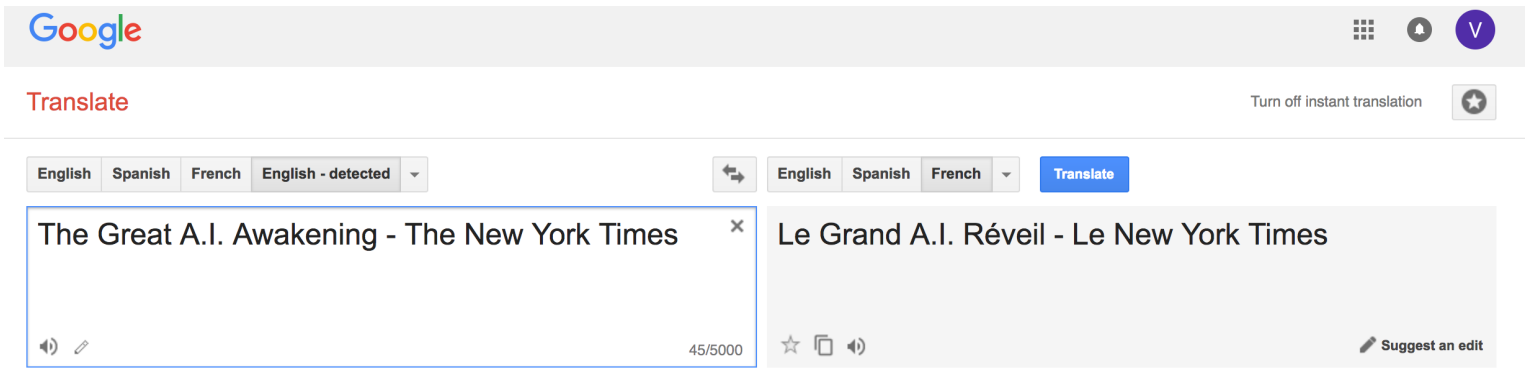

Turing's Test [1]

- Can machines think?
- Simple version - Can machines imitate humans?
- E.g., ELIZA
  - `http://www.manifestation.com/neurotoys/eliza.php3`
- Difference between "human behavior" and "intelligent behavior".

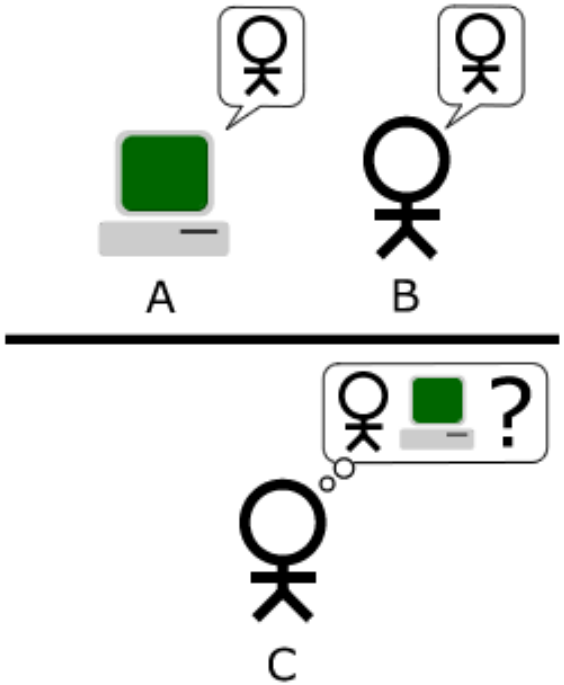

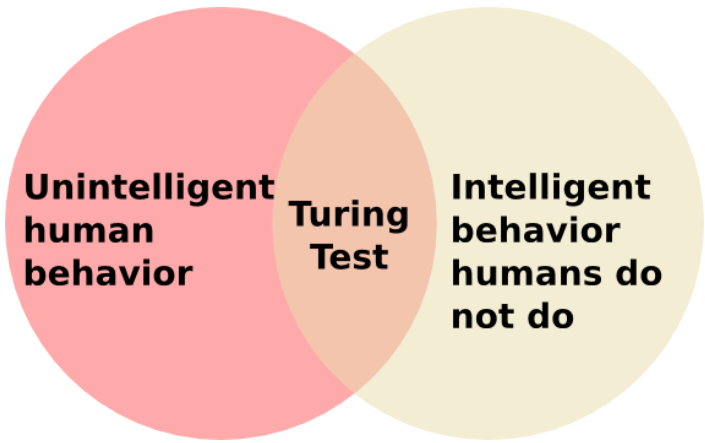

What makes an artificial machine **intelligent**?

1. Talk. See. Hear.
   - Natural Language Processing, Computer Vision, Speech Recognition
2. Store. Access. Represent. (Knowledge)
   - Ontologies. Semantic Networks. Information Retrieval.
3. Reason.
   - Mathematical Logic. Bayesian Inference.
4. **Learn.**
   - Improve with Experience

What makes an artificial machine **intelligent**?

1. Talk. See. Hear.
   - Natural Language Processing, Computer Vision, Speech Recognition
2. Store. Access. Represent. (Knowledge)
   - Ontologies. Semantic Networks. Information Retrieval.
3. Reason.
   - Mathematical Logic. Bayesian Inference.
4. **Learn.**
   - Improve with Experience
     - Machine Learning

Human Learning?

- What do we learn?
  - Concepts (this is a chair, that is not a chair)
  - Distinguishing concepts (this is a chair, that is a table)
  - Other things (language, juggling, using a remote)
- How do we learn?
  1. Teaching (Passive).
  2. Experience (Active).
     - 2.1 Examples.
     - 2.2 Queries.
     - 2.3 Experimentation.

What is Machine Learning?

- Computers learn without being **explicitly programmed**.
  - Arthur Samuel (1959)
- A computer program learns from experience E with respect to some task T, if its performance P while performing task T improves over E.
  - Tom Mitchell (1989)

Why Machine Learning?

- Machines that know everything from the beginning?
  - Too bulky. Creator already knows everything. Fails with new experiences.
- Machines that learn?
  - Compact. Learn what is necessary.
  - Adapt.
  - Assumption: Future experiences are not too different from past experiences.
    - Have (structural) relationship.

Example - A Robotic Furniture Sorter

- Our future overlords (Kiva)
  - https://youtube.googleapis.com/v/6KRjuuEVEZs

Make Kiva Learn

- Distinguish between **tables** and **chairs**.
- How do you represent the input?
- Distinguishing features.
  - has back rest?
- Good features should be:
  - Computable. Distinguishing. Representative.
  - is a chair?
  - has four legs?
  - is white?

Learning from (Past) Data

Deductive Logic

- All birds can fly
- Dodo is a bird
- $\Rightarrow$ Dodo can fly

Inductive Logic

- A stingray can swim
- Stingray is a fish
- $\Rightarrow$ All fish can swim

Learning from (Past) Data

Deductive Logic

- All birds can fly
- Dodo is a bird
- $\Rightarrow$ Dodo can fly

Inductive Logic

- A stingray can swim
- Stingray is a fish
- $\Rightarrow$ All fish can swim

Core Tenet

- ~~Deduce~~ Induce from past
- Generalize for future

A Simple ML Example

https://quickdraw.withgoogle.com/

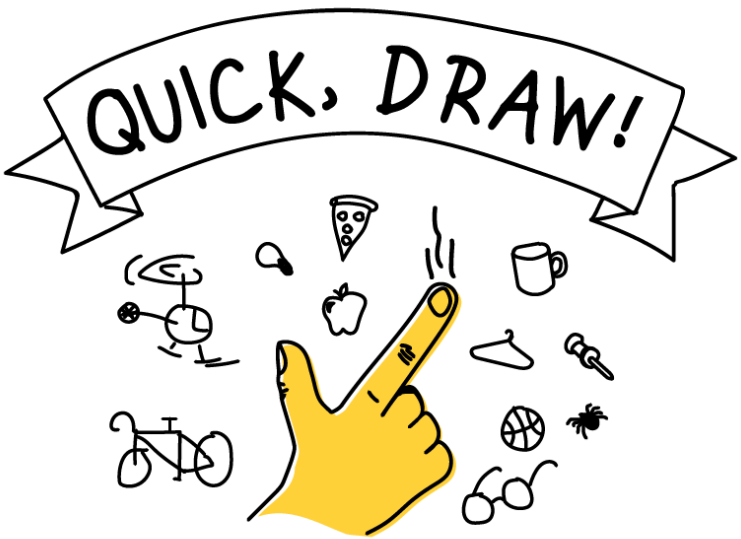

References


A. M. Turing.
Computing machinery and intelligence.
Mind, 59:433–460, 1950.